

Airflow on Kubernetes

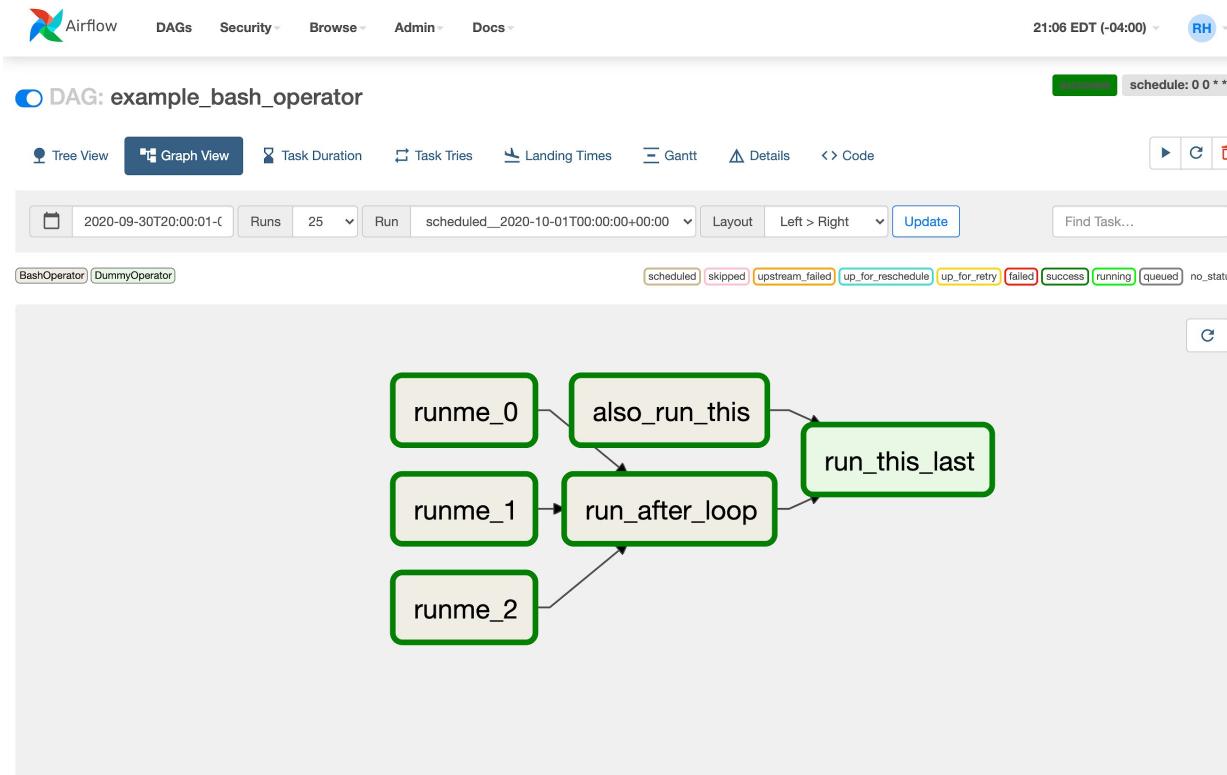
tips, tricks and pitfalls

| What is airflow?

- ▼ Workflow
 - ▽ Execution
 - ▽ Scheduling
 - ▽ Monitoring
- ▼ Extensible (providers, plugin)
- ▼ Scalable
- ▼ Python



Vocabulary

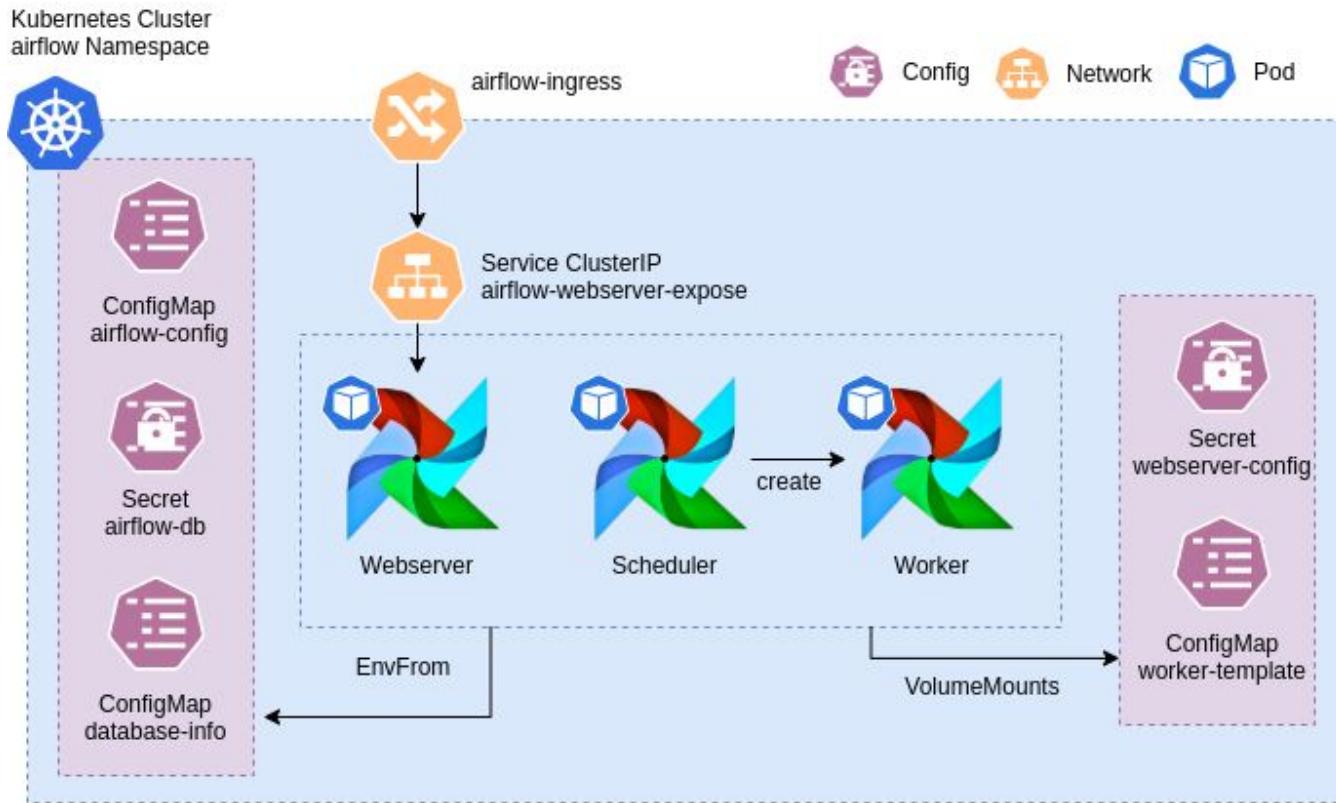


| What is Kubernetes?

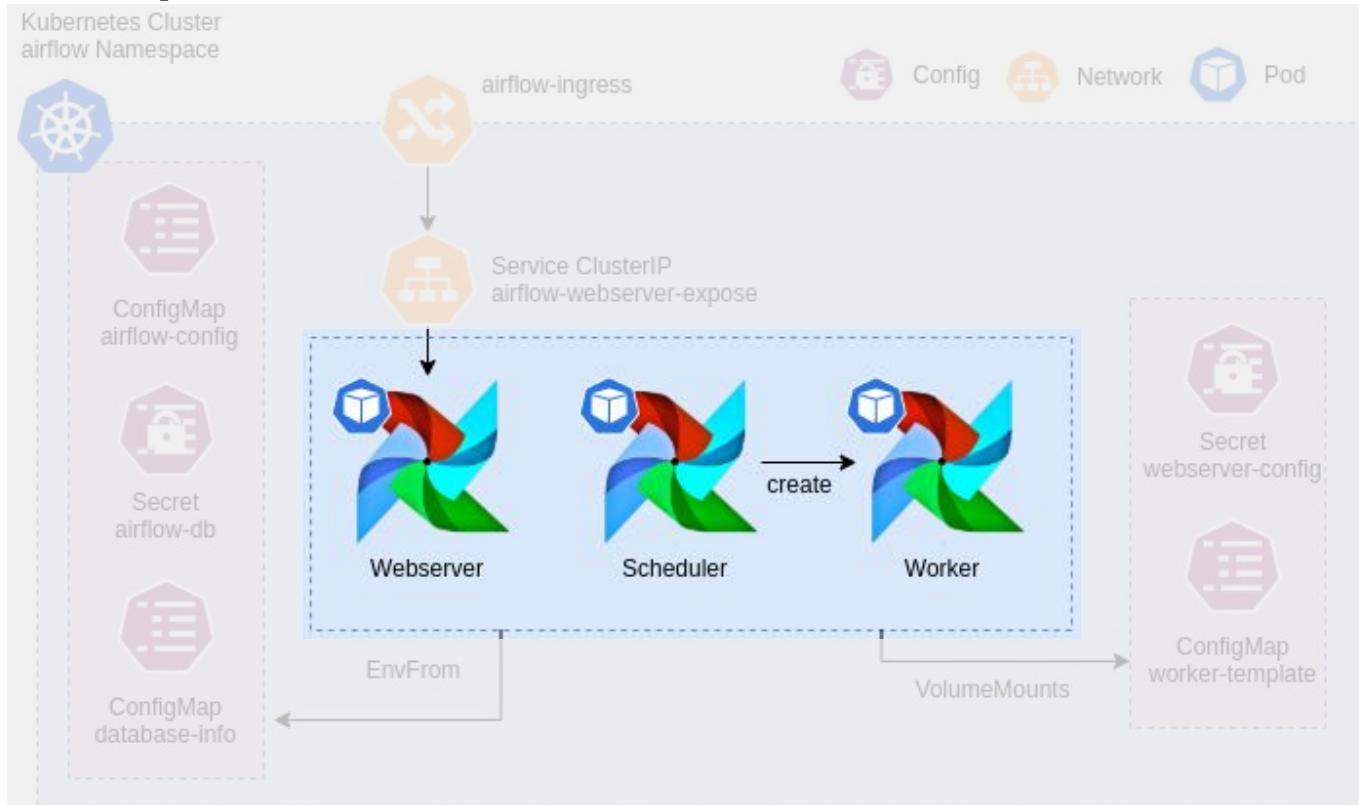
- ▼ Containerized workloads
 - ▼ Manage
 - ▼ Self Heal
 - ▼ Scale
- ▼ Scalable
- ▼ High Availability
- ▼ Go



Airflow on K8s Overview



Airflow Components



Webserver

The screenshot shows the Airflow Webserver interface with the title "DAGs". At the top, there are navigation links: Airflow logo, DAGs, Security, Browse, Admin, and Docs. On the right, it shows the time "21:11 UTC" and a user icon labeled "RH". Below the header, there are three buttons: All (26), Active (10), and Paused (16). There are also two search bars: "Filter DAGs by tag" and "Search DAGs". The main table lists 10 DAGs from the "example" package:

DAG	Owner	Runs	Schedule	Last Run	Recent Tasks	Actions	Links
example_bash_operator example example2	airflow	2	0 * * * *	2020-10-26, 21:08:11	6		
example_branch_dop_operator_v3 example	airflow	0	* /1 * * *		0		
example_branch_operator example example2	airflow	1	@daily	2020-10-23, 14:09:17	11		
example_complex example example2 example3	airflow	1	None	2020-10-26, 21:08:04	37		
example_external_task_marker_child example	airflow	1	None	2020-10-26, 21:07:33	2		
example_external_task_marker_parent example	airflow	1	None	2020-10-26, 21:08:34	1		
example_kubernetes_executor example example2	airflow	0	None		0		
example_kubernetes_executor_config example3	airflow	1	None	2020-10-26, 21:07:40	5		
example_nested_branch_dag example	airflow	1	@daily	2020-10-26, 21:07:37	9		
example_passing_params_via_test_command example	airflow	0	* /1 * * *		0		

WebUI

REST API

Authentication

| Scheduler

- ▼ Parse DAG
- ▼ Trigger DAG and Task
- ▼ Monitor DAG and Task
- ▼ High Availability since airflow 2.0
- ▼ Check Concurrency and resource Limits

Workers

- ▼ Execute the Task
- ▼ Can use different backend
 - ▽ Dask
 - ▽ Celery
 - ▽ Kubernetes
 - ▽ Local

| DB

- ▼ Configurable
 - ▽ MySQL
 - ▽ Postgres
 - ▽ SQLite (for dev only)
- ▼ Airflow use SQLAlchemy
- ▼ Used by scheduler, webserver and worker
- ▼ Store DAG and Task Execution
- ▼ Store DAG and Task State
- ▼ Store xcom

Let's Begin

Webserver

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: airflow-webserver
spec:
  replicas: 1
  revisionHistoryLimit: 1
  selector:
    matchLabels:
      app: airflow-webserver
  template:
    metadata:
      labels:
        app: airflow-webserver
    spec:
      containers:
        - image: apache/airflow:2.0.0-python3.8
          name: airflow-webserver
          command: ["/entrypoint", "webserver"]
```

Scheduler

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: airflow-scheduler
spec:
  replicas: 1
  revisionHistoryLimit: 1
  selector:
    matchLabels:
      app: airflow-scheduler
  template:
    metadata:
      labels:
        app: airflow-scheduler
    spec:
      containers:
        - image: apache/airflow:2.0.0-python3.8
          name: airflow-scheduler
          command: ["/entrypoint", "scheduler"]
```

The End

```
sqlalchemy.exc.OperationalError: (sqlite3.OperationalError) no such table: dag
```

Init DB

- ▼ `!`use volume and set pgdata``
- ▼ separate database info from secret since you can use the database for other application
- ▼ if you reuse this DB you will have to create a database with correct user and permission before launching airflow
- ▼ if airflow don't find database on startup it fail

```
apiVersion: v1
kind: Secret
metadata:
  name: airflow-db
type: Opaque
stringData:
  POSTGRES_USER: airflow
  POSTGRES_PASSWORD: airflow
  POSTGRES_DB: airflow
  FERNET_KEY: "cxIPNj14G6trLTpzFvVpPdi1yZWnnbsnzzKanFyWfY="
```

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: db
    name: "postgres-db"
spec:
  ports:
    - port: 5432
      protocol: TCP
      targetPort: 5432
  selector:
    app: db
  type: ClusterIP
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: database-info
data:
  POSTGRES_HOST: postgres-db
  POSTGRES_PORT: "5432"
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: db
    name: postgres-db
spec:
  replicas: 1
  revisionHistoryLimit: 1
  selector:
    matchLabels:
      app: db
  template:
    metadata:
      labels:
        app: db
    spec:
      containers:
        - image: postgres:13.0-alpine
          name: postgres
          envFrom:
            - secretRef:
                name: airflow-db
            - configMapRef:
                name: database-info
  ports:
    - containerPort: 5432
```

FERNET_KEY: {{ randAscii 32 | b64enc | quote }}

| Set DB parameter in Airflow

- ▼ *envFrom* ConfigMap and secret are populated before *env*
- ▼ set DB env variable in webserver scheduler and worker
- ▼ Fernet key is used to encrypt connection password and variables

```
envFrom:  
- secretRef:  
  name: airflow-db  
- configMapRef:  
  name: database-info  
env:  
- name: AIRFLOW__CORE__SQL_ALCHEMY_CONN  
  value: "postgres://$(POSTGRES_USER):$(POSTGRES_PASSWORD)@$(POSTGRES_HOST):$(POSTGRES_PORT)/$(POSTGRES_DB)"  
- name: AIRFLOW__CORE__FERNET_KEY  
  value: "$(FERNET_KEY)"
```

```
sqlalchemy.exc.OperationalError: (sqlite3.OperationalError) no such table: dag
```

Init Container

- ▼ You have to init Airflow db with init container
- ▼ Always use *db upgrade* it will create tables or upgrade them
- ▼ init DB in only one airflow container (webserver)
- ▼ the other component will start and fail during tables creation

```
initContainers:  
- name: upgrade-db  
  image: apache/airflow:2.0.0-python3.8  
  command: [ "/entrypoint", "db", "upgrade" ]  
  envFrom:  
    - secretRef:  
        name: airflow-db  
    - configMapRef:  
        name: database-info  
  env:  
    - name: AIRFLOW__CORE__SQLALCHEMY_CONN  
      value: "postgres://$(POSTGRES_USER):$(POSTGRES_PASSWORD)@$(POSTGRES_HOST):$(POSTGRES_PORT)/$(POSTGRES_DB)"
```


The End

WebUI ?

Service

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: airflow-webserver
  name: airflow-expose
spec:
  ports:
    - port: 8080
      protocol: TCP
      targetPort: 8080
  selector:
    app: airflow-webserver
  type: ClusterIP
```

Ingress

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: airflow-ingress
  annotations:
    kubernetes.io/ingress.class: nginx
spec:
  rules:
    - host: airflow.local
      http:
        paths:
          - path: /airflow
            pathType: Prefix
            backend:
              serviceName: airflow-expose
              servicePort: 8080
```

Config

- ▼ Airflow can use environment variable as config
- ▼ You can also set config with airflow.cfg file
- ▼ ConfigMap will ensure airflow params consistency across service (webserver, scheduler)
- ▼ Using env is easier than putting airflow.cfg in a ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: airflow-config
data:
  AIRFLOW__WEBSERVER__BASE_URL: "http://minikube.info/airflow"
  AIRFLOW__WEBSERVER__ENABLE_PROXY_FIX: "True"
  AIRFLOW__WEBSERVER__EXPOSE_CONFIG: "True"
```

```
envFrom:
- configMapRef:
    name: airflow-config
```



Airflow

10:18 UTC

Log In

Access is Denied



Sign In

Enter your login and password below:

Username:

Password:

Sign In

[2021-01-20 10:19:03,988] {manager.py:727} WARNING - No user yet created, use flask fab command to do it.

```
$ airflow users create --username admin --role Admin --password admin \
--firstname FIRST_NAME --lastname LAST_NAME --email admin@example.org
```

The End

DAGs

| Where is my DAG?

- ▼ /opt/airflow/dags (`AIRFLOW_CORE_DAGS_FOLDER`)
- ▼ mandatory in scheduler
- ▼ in webserver if you want to see code in the UI
- ▼ using configMap mounted in container doesn't work

Where is my DAG?

	In Image	Git Sync	Shared Volume
pros	easy setup no poll DAG consistency no volume	RWO Volume	DAG consistency across service Can use a webserver or something else to push DAG in the volume
cons	build and deploy an entire airflow for each DAG update	poll delay pull an entire repository (unless you use some git black magic with a custom git-sync container) DAG inconsistency Deploy key or public repository	need RWX volumes

Git Sync

- ▼ use kubernetes [git-sync](#)
- ▼ Create a multi container pod
- ▼ Share an emptyDir volume
- ▼ Use an init container for Worker
- ▼ Share ssh keys with kubernetes secrets

```
containers:  
- image: k8s.gcr.io/git-sync/git-sync:v3.2.2  
  name: git-sync  
  env:  
    - name: GIT_SYNC_REPO  
      value: git@github.com:teamclairvoyant/airflow-maintenance-dags.git  
    - name: GIT_SYNC_BRANCH  
      value: airflow2.0  
    - name: GIT_SYNC_WAIT  
      value: "60"  
    - name: GIT_SYNC_SSH  
      value: "true"  
    - name: GIT_SYNC_ROOT  
      value: /tmp/git/  
    - name: GIT_SSH_KEY_FILE  
      value: /tmp/ssh/airflow-id-rsa  
    - name: GIT_KNOWN_HOSTS  
      value: "false"  
    - name: GIT_SYNC_ADD_USER  
      value: "true"  
  volumeMounts:  
    - name: dags  
      mountPath: /tmp/git  
    - name: airflow-git-sync-ssh  
      mountPath: /tmp/ssh
```

Kubernetes Executor

- ▼ set worker pod template in `/opt/airflow/template` folder
- ▼ use security context to mount this volume with the right owner (50000)
- ▼ create a custom image with kubernetes dependencies
- ▼ configure worker image in airflow.cfg or with env var
- ▼ configure RBAC to allow the scheduler to create worker pod
- ▼ configure RBAC to allow the webserver to get pods logs
- ▼ configure RBAC to allow worker to create pod if you want to use KubernetesPodOperator

```
securityContext:  
  runAsNonRoot: true  
  runAsUser: 50000  
  runAsGroup: 50000  
  fsGroup: 50000
```

Worker Pod Template

	in Image	in ConfigMap
pros	easy setup	hot editing deploy with other airflow services
cons	no hot editing	

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: worker-pod-template
data:
  worker_pod_template.yaml: |-
    apiVersion: v1
    kind: Pod
    metadata:
      name: airflow-worker
    spec:
      serviceAccountName: airflow-worker
      containers:
        - image: apache/airflow:2.0.0-python3.8
          name: base
          envFrom:
            - secretRef:
                name: airflow-db
            - configMapRef:
                name: database-info
            - configMapRef:
                name: airflow-config
      env:
        - name: AIRFLOW__CORE__EXECUTOR
          value: LocalExecutor
        - name: AIRFLOW__CORE__SQLALCHEMY_CONN
          value: "postgres://..."
        - name: AIRFLOW__CORE__FERNET_KEY
          value: "$(FERNET_KEY)"
    restartPolicy: Never
    securityContext:
      runAsNonRoot: true
      runAsUser: 50000
      runAsGroup: 50000
      fsGroup: 50000
```

KubernetesExecutor config

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: airflow-config
data:
  ...
  AIRFLOW__CORE__EXECUTOR: "KubernetesExecutor"
  AIRFLOW__KUBERNETES__NAMESPACE: "airflow"
  AIRFLOW__KUBERNETES__POD_TEMPLATE_FILE: "/opt/airflow/template/worker_pod_template.yaml"
  AIRFLOW__KUBERNETES__WORKER_SERVICE_ACCOUNT_NAME: "airflow-worker"
  AIRFLOW__KUBERNETES__WORKER_CONTAINER_REPOSITORY: "apache-airflow"
  AIRFLOW__KUBERNETES__WORKER_CONTAINER_TAG: "2.0.0-python3.8"
  AIRFLOW__KUBERNETES__DELETE_WORKER_PODS: "True"
  AIRFLOW__KUBERNETES__DELETE_WORKER_PODS_ON_FAILURE: "True"
  AIRFLOW__KUBERNETES__IN_CLUSTER: "True"
```

RBAC: scheduler

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: airflow-scheduler-k8s-auth
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list", "watch", "create", "patch", "delete"]
- apiGroups: [""]
  resources: ["pods/log"]
  verbs: ["get"]
- apiGroups: [""]
  resources: ["pods/exec"]
  verbs: ["create", "get"]
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: airflow-scheduler-k8s-auth
subjects:
- kind: ServiceAccount
  name: airflow-scheduler
roleRef:
  kind: Role
  name: airflow-scheduler-k8s-auth
  apiGroup: rbac.authorization.k8s.io
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    name: airflow-scheduler
  name: airflow-scheduler
```

RBAC: webserver

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: airflow-webserver-k8s-auth
rules:
  - apiGroups: [""]
    resources: ["pods/log"]
    verbs: ["get"]
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: airflow-webserver-k8s-auth
subjects:
  - kind: ServiceAccount
    name: airflow-webserver
roleRef:
  kind: Role
  name: airflow-webserver-k8s-auth
  apiGroup: rbac.authorization.k8s.io
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    name: airflow-webserver
  name: airflow-webserver
```

RBAC: worker

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: airflow-worker
rules:
- apiGroups: []
  resources: ["pods"]
  verbs: ["get", "list", "watch", "create", "patch", "delete"]
- apiGroups: []
  resources: ["pods/log"]
  verbs: ["get"]
- apiGroups: []
  resources: ["pods/exec"]
  verbs: ["create", "get"]
- apiGroups: []
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: airflow-worker
subjects:
- kind: ServiceAccount
  name: airflow-worker
roleRef:
  kind: Role
  name: airflow-worker
  apiGroup: rbac.authorization.k8s.io
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    name: airflow-worker
  name: airflow-worker
```

| Build your own Image

- ▼ why?

- ▽ dag in image

- ▽ custom providers (aws, k8s)

- ▼ --user because airflow image are not root

- ▼ --chown when you copy dags

```
FROM apache/airflow:2.0.0-python3.6

RUN pip install --user --no-cache-dir \
    apache-airflow-providers-amazon \
    apache-airflow-providers-cncf-kubernetes

COPY --chown=airflow:airflow dags /opt/airflow/dags/
```

| Where are my Logs?

- ▼ Logs are in kubernetes worker that get destroy when task end
- ▼ You have to enable **AIRFLOW__LOGGING__REMOTE_LOGGING**
- ▼ Use service account to allow worker to write on the logs bucket

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: airflow-config
data:
  ...
  AIRFLOW__LOGGING__REMOTE_LOGGING: "True"
  AIRFLOW__LOGGING__REMOTE_BASE_LOG_FOLDER: "s3://logs-airflow"
```

Where are my Logs? Azure Edition

- ▼ You need to create a wasb connection on airflow (extra can take a `connection_string`)
- ▼ You need to define a `log_config.py` (example is [airflow local settings.py](#))
- ▼ override `wasb_container` with your container name or fetch it from config
- ▼ add the `log_config.py` in `PYTHONPATH`
- ▼ use “[`secret backend`](#)” to set connection

```
if REMOTE_LOGGING:  
    REMOTE_BASE_LOG_FOLDER: str = conf.get('logging', 'REMOTE_BASE_LOG_FOLDER')  
    ...  
    elif REMOTE_BASE_LOG_FOLDER.startswith('wasb'):  
        WASB_CONTAINER: str = conf.get('logging', 'WASB_CONTAINER')  
        WASB_REMOTE_HANDLERS: Dict[str, Dict[str, Union[str, bool]]] = {  
            'task': {  
                'class': 'airflow.providers.microsoft.azure.log.wasb_task_handler.WasbTaskHandler',  
                'formatter': 'airflow',  
                'base_log_folder': str(os.path.expanduser(BASE_LOG_FOLDER)),  
                'wasb_log_folder': REMOTE_BASE_LOG_FOLDER,  
                'wasb_container': WASB_CONTAINER,  
                'filename_template': FILENAME_TEMPLATE,  
                'delete_local_copy': False,  
            },  
        }  
  
        DEFAULT_LOGGING_CONFIG['handlers'].update(WASB_REMOTE_HANDLERS)
```

| Auth (LDAP)

- ▼ LDAP and other auth provider are configured in /opt/airflow/webserver_config.py
- ▼ it's use Flask [appbuilder](#) and [flask-LDAP](#)
- ▼ you have to install flask LDAP dependency in airflow image
 - ▽ apt-get install -y build-essential gcc libsasl2-dev libldap2-dev libssl-dev
 - ▽ pip install --user python-ldap
- ▼ store webserver_config.py in secret and use subpath to mount it

```
volumeMounts:  
- name: airflow-webserver-config  
  mountPath: /opt/airflow/webserver_config.py  
  subPath: config.py
```

```
volumes:  
- name: airflow-webserver-config  
  secret:  
    secretName: airflow-webserver-config
```

Liveness

▼ Webserver

▽ Liveness on /health

▽ /health always return 200 even if DB or Scheduler are down

▼ Scheduler

▽ Custom heartbeat dag

```
livenessProbe:  
  httpGet:  
    path: /health  
    port: 8080  
  periodSeconds: 600  
  initialDelaySeconds: 240
```

Scheduler Liveness

- ▼ create a DAG that update a timestamp in DB
- ▼ check that last timestamp in DB is less than delay
- ▼ we can't rely on airflow healthcheck
 - ▼ it's depends on webserver
 - ▼ it don't test that the scheduler can schedule dag
- ▼ It's recommended by [airflow documentation](#)
- ▼ may fail if you reach pool limit

```
#!/usr/local/bin/python
from sqlalchemy.exc import ProgrammingError
from airflow import settings

session = settings.Session()
session.execute("INSERT INTO healthcheck(worker_heartbeat_ts) VALUES (current_timestamp);")
session.commit()
session.close()
```

```
from datetime import datetime, timedelta
from airflow.models import DAG
from airflow.operators.bash import BashOperator

"""
Updates timestamp in the airflow db, healthcheck table
"""

default_args = {
    'owner': 'operations',
    'depends_on_past': False,
    'execution_timeout': timedelta(minutes=1),
}

dag = DAG('zz-airflow-heartbeat-dag',
          default_args=default_args,
          schedule_interval=timedelta(minutes=10),
          max_active_runs=1,
          start_date=(datetime.now() - timedelta(days=1)),
          is_paused_upon_creation=False,
          catchup=False)

update_task = BashOperator(
    task_id='update_timestamp',
    bash_command="/usr/local/bin/python /opt/airflow/bin/heartbeat.py",
    dag=dag
)
```

Scheduler Liveness

```
livenessProbe:  
  exec:  
    command:  
      - python  
      - /opt/airflow/bin/healthcheck.py  
failureThreshold: 1  
initialDelaySeconds: 600  
periodSeconds: 1200  
successThreshold: 1  
timeoutSeconds: 20
```

```
#!/usr/local/bin/python  
import sys  
  
from airflow import settings  
from datetime import datetime, timezone  
  
# 20 minutes, the maximum number of seconds that we wait for the liveness task to update value in  
the database.  
# After this number is reached, kubernetes will kill the scheduler pod.  
max_execution_delay = 20 * 60  
session = settings.Session()  
  
result = session.execute("select max(worker_heartbeat_ts) from healthcheck").fetchall()  
session.commit()  
session.close()  
  
if (len(result)) == 0:  
    raise Exception('No execution_ts found!')  
else:  
    most_recent_ts = dict(result[0])['max']  
  
delta = datetime.now(timezone.utc) - most_recent_ts  
print(f'Last execution was {delta.seconds} seconds ago')  
  
if delta.seconds > max_execution_delay: # 20 minutes  
    raise Exception(f'Last execution was on {most_recent_ts}. This is {delta.seconds} second ago,  
maximum permitted: {max_execution_delay}!')  
  
print("everything is fine")  
sys.exit(0)
```

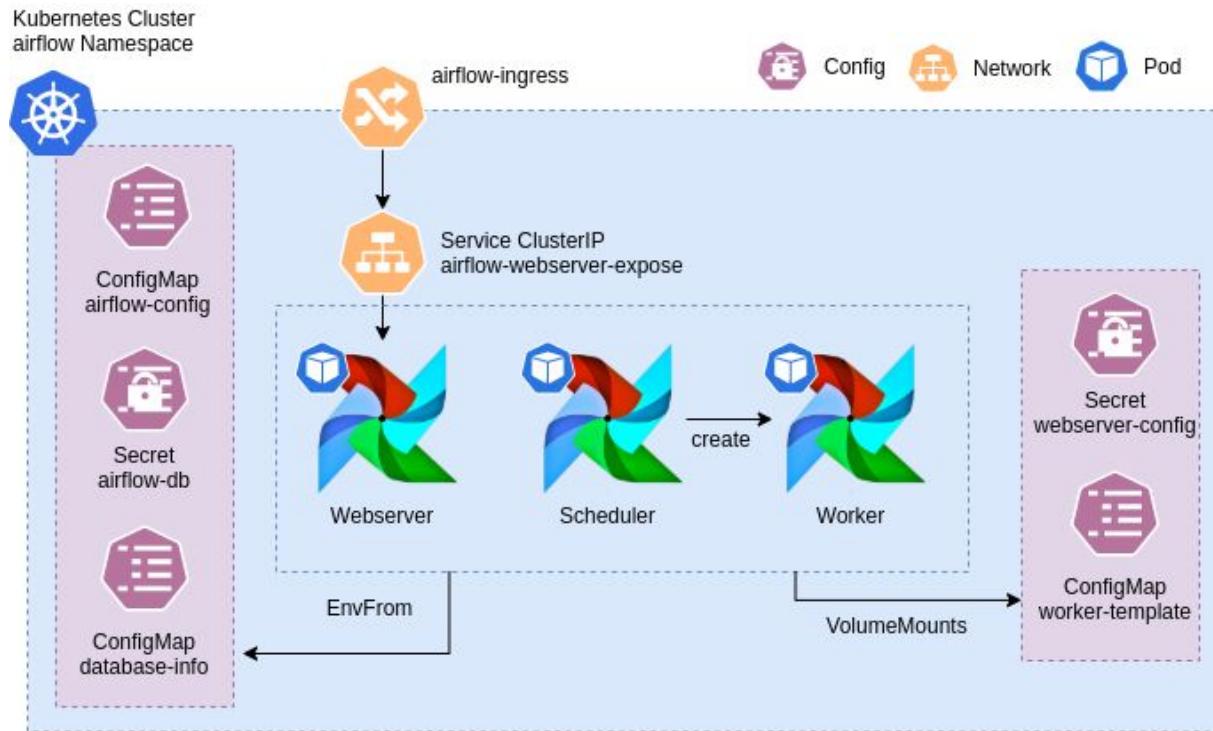
Scheduler Liveness (2)

- ▼ You can use scheduler heartbeat
- ▼ Don't fail if you reach pool limit
- ▼ May not fail when scheduler deadlock

```
livenessProbe:  
  initialDelaySeconds: 600  
  timeoutSeconds: 5  
  failureThreshold: 10  
  periodSeconds: 30  
  exec:  
    command:  
      - python3  
      - /opt/airflow/bin/healthcheck.py
```

```
#!/usr/local/bin/python  
from airflow.jobs.scheduler_job import SchedulerJob  
from airflow.utils.db import create_session  
from airflow.utils.net import get_hostname  
import sys  
with create_session() as session:  
    job = session.query(SchedulerJob).filter_by(hostname=get_hostname()).order_by(SchedulerJob.latest_heartbeat.desc()).limit(1).first()  
    sys.exit(0 if job.is_alive() else 1)
```

Airflow on K8s Overview



| One More Thing

- ▼ use one navbar color by env `AIRFLOW__WEB SERVER__NAVBAR_COLOR`
- ▼ Use podAntiAffinity to avoid having scheduler and webserver on the same node if they are small (2 Core 4Go)
- ▼ Use Dedicated Nodes for airflow and other static workload (taint or labels with nodeSelector)
- ▼ Use autoscalling for airflow worker (taint or labels with nodeSelector)
- ▼ You must have a decent machine for your database (2core)

Generate Random secret on helm

```
/* 1. Query your secret for "existence" and return in the $secret variable*/
{{ $secret_airflow := (lookup "v1" "Secret" .Release.Namespace "airflow-db") -}}
apiVersion: v1
kind: Secret
metadata:
  name: airflow-db
type: Opaque
/* 2. If the secret exists, write it back out (it'll be Base64 encoded so used the "data" key)*/
{{ if $secret_airflow -}}
data:
  POSTGRES_USER: {{ $secret_airflow.data.POSTGRES_USER }}
  POSTGRES_PASSWORD: {{ $secret_airflow.data.POSTGRES_PASSWORD }}
  POSTGRES_DB: {{ $secret_airflow.data.POSTGRES_DB }}
  FERNET_KEY: {{ $secret_airflow.data.FERNET_KEY }}
/* 3. If it doesn't exist ... create it (this time as "stringData" as it will be a raw value) !*/
{{ else -}}
stringData:
  POSTGRES_USER: airflow
  POSTGRES_PASSWORD: {{ randAlphaNum 32 | quote }}
  POSTGRES_DB: airflow
  FERNET_KEY: {{ randAscii 32 | b64enc | quote }}
{{ end }}
```

The End

for real
Any Questions?