


Deeplearning in production

the Data Engineer part

whoami



Scauglog
Data Engineer, Xebia

<https://github.com/scauglog/prez>

init project


| Team Astro



Product Owner



Scrum Master

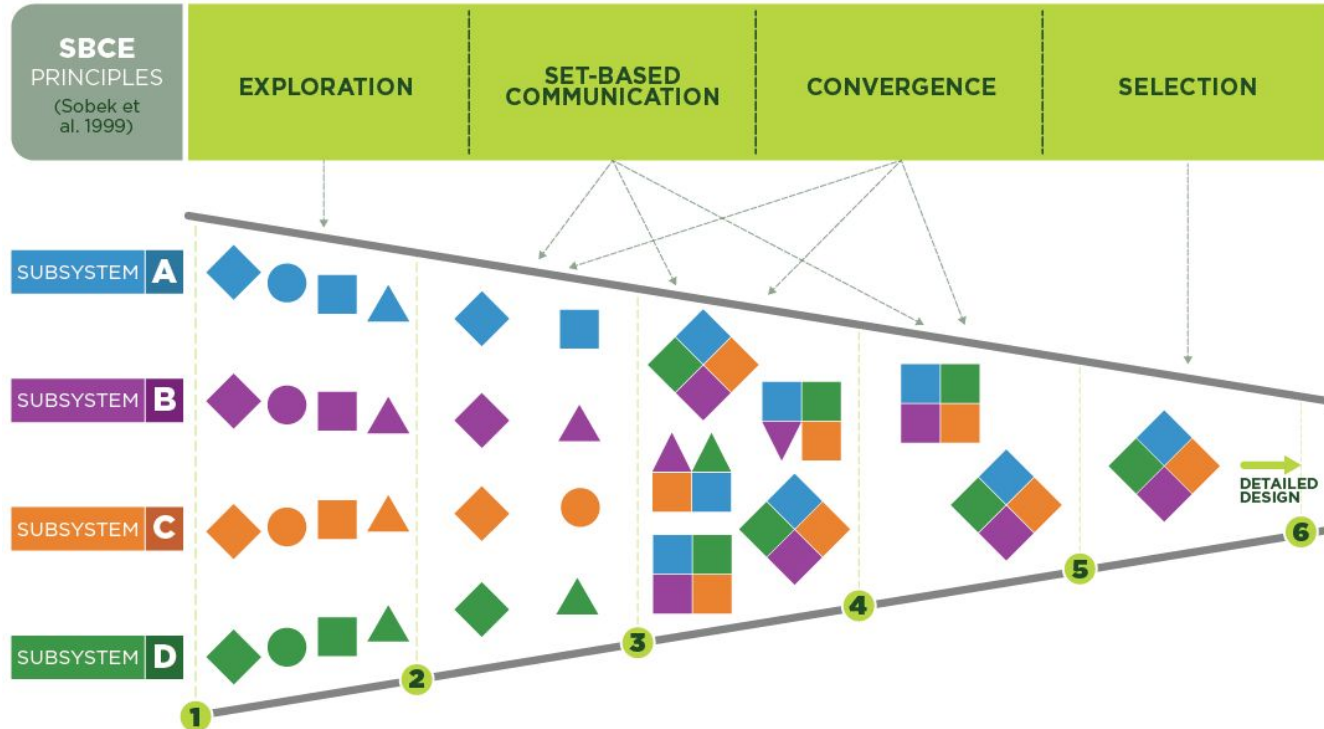


Data Scientists, Data Engineers,
Machine Learning Engineers

| Business

- ▼ Buy sponsored link on google adwords
- ▼ 10M Predictions in less than 1 hour (~2700/s)
- ▼ Bid each day
- ▼ Each bid should cost less than what we earn

Choose your model




| And the winner is


XGBoost



| And the winner is



| What is Deep Learning?



| What is a Deep Learning model?



protobuf
Protocol Buffers


| Deeplearning in Production at Scale





Choose your Framework

- ▼ Distributed Prediction
- ▼ Can create complex network
- ▼ Documentation
- ▼ Community




DL4J




| And the winner is



| Wait




| And the winner is




DL4J

How to Deep Learn

Train Workflow



Predict Workflow






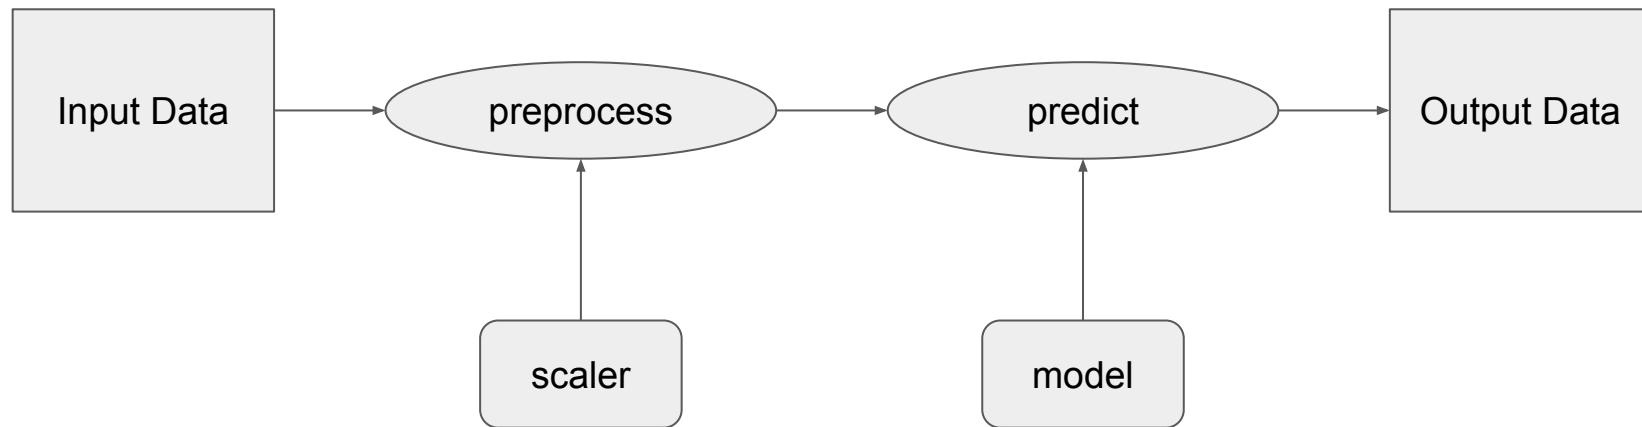
| Preprocessing

- ▼ Scaling (normalisation, min max, ...)
- ▼ Replace null
- ▼ Lagging

Train Workflow



Predict Workflow



Predict at scale

Scaling Prediction: naïve approach

```
def predict(input: RDD[PreprocessRow], modelPath: String): RDD[SinglePredictionRow]
= {
  input.map { row =>
    // Load model
    val hdfs = FileSystem.get(new Configuration())
    val source = new Path(modelPath)
    val model = ModelSerializer.restoreComputationGraph(hdfs.open(source), true)

    // make prediction
    val prediction = model.output(row.features)(0).getColumn(0).toFloatVector

    // return prediction
    SinglePredictionRow.fromPreprocessRow(row, prediction(0))
  }
}
```

Scaling Prediction: faster

```
def predict(input: RDD[PreprocessRow], modelPath: String): RDD[SinglePredictionRow]
= {
  // Load model
  lazy val hdfs = FileSystem.get(new Configuration())
  lazy val source = new Path(modelPath)
  lazy val model = ModelSerializer.restoreComputationGraph(hdfs.open(source), true)

  input.map { row =>
    // make prediction
    val prediction = model.output(row.features)(0).getColumn(0).toFloatVector

    // return prediction
    SinglePredictionRow.fromPreprocessRow(row, prediction(0))
  }
}
```

Scaling Prediction: fastest

```
def predict(input: RDD[PreprocessRow], modelPath: String):  
  RDD[SinglePredictionRow] = {  
    // Load model  
    lazy val hdfs = FileSystem.get(new Configuration())  
    lazy val source = new Path(modelPath)  
    lazy val model = ModelSerializer.restoreComputationGraph(hdfs.open(source),  
true)  
  
    input.mapPartitions { partition =>  
      val partitionSeq = partition.toSeq  
      if (partitionSeq.isEmpty) {  
        Iterator(): Iterator[SinglePredictionRow]  
      } else {  
        val features = partitionSeq.map(_.features).reduce( (x, y) => Nd4j.concat(0,  
x, y))  
  
        val predictions = model.output(features)(0).getColumn(0).toFloatVector  
        partitionSeq.zip(predictions).map { case (row, prediction) =>  
          SinglePredictionRow.fromPreprocessRow(row, prediction)  
        }.toIterator  
      }  
    }  
  }
```

OOM

- ▼ ND4J Array are C++ offheap object
- ▼ Cache your dataframe or look stage details to estimate memory usage
- ▼ Set *spark.yarn.executor.memoryOverhead*
- ▼ Use ND4J workspace to properly manage memory deallocation
- ▼ Repartition your dataframe before prediction to ensure equals partition
- ▼ Set *spark.sql.shuffle.partitions*

OOM

```
def predict(input: RDD[PreprocessRow], modelPath: String, numFeatures: Int, timeSteps: Int): RDD[SinglePredictionRow] =
{
  // Load model ...
  lazy val basicConfig: WorkspaceConfiguration = WorkspaceConfiguration.builder().initialSize(0)
    .policyLearning(LearningPolicy.NONE).policyAllocation(AllocationPolicy.STRICT).build()
  lazy val workspace = Nd4j.getWorkspaceManager.getAndActivateWorkspace(basicConfig, "myWorkspace")

  input.mapPartitions { partition =>
    val partitionSeq = partition.toSeq
    if (partitionSeq.isEmpty) {
      Iterator(): Iterator[SinglePredictionRow]
    } else {
      workspace.notifyScopeEntered()
      val features = Nd4j.create(partitionSeq.flatMap(_.features).toArray, Array(partitionSeq.size, numFeatures,
timeSteps))
      val predictions = model.output(false, workspace, features)(0).toFloatVector
      workspace.notifyScopeLeft()

      partitionSeq.zip(predictions).map { case (row, prediction) =>
        SinglePredictionRow.fromPreprocessRow(row, prediction)
      }.toIterator
    }
  }
}
```

| Compile

- ▼ Maven
- ▼ `-Djavacpp.platform=linux-x86_64`
- ▼ Exclude
 - ▼ `deeplearning4j-datasets`
 - ▼ `deeplearning4j-datavec-iterators`
 - ▼ `deeplearning4j-ui-components`




Training at scale


| Retrain again and again and again...

- ▼ Model performance decline over time
- ▼ Hyperparameter tuning
- ▼ Deep Learning model rarely comes alone (clustering)

Predict Workflow



Train Workflow




| Training at scale: AWS EC2



Training at scale: AWS EC2

- ▼ Create VPC
- ▼ Create Subnet associated to VPC
- ▼ Create an IGW associated to VPC
- ▼ Create a route table associated to IGW
- ▼ Create a Security Group associated to VPC
 - ▽ Authorize ssh only for my IP
- ▼ Create a key pair
- ▼ Create EC2 server with EBS volume



| Training at scale: AWS EC2

- ▼ Add ssh keys of team members
- ▼ Install cuda, cudnn, nccl and configure them
- ▼ Deploy train jar to EC2 instance
- ▼ Deploy train pipeline to EC2 instance
- ▼ Deploy preprocessed data to EC2 instance
- ▼ Deploy auto shutdown script

| Training at scale: AWS EC2


- ▼ Ansible
- ▼ Transfert preprocess data to S3
- ▼ Store model in S3
- ▼ Check CPU vs GPU training time
- ▼ Keep track of training config and performance
- ▼ Share knowledge with Data Scientist
- ▼ Put your data in EBS volume if they fit



ANSIBLE

Training with DL4J: Lessons learned

- ▼ Beware of tensor order
- ▼ Prefetch data in memory (InMemoryDatasetIterator)
- ▼ Add listener to monitor your training compute performance
- ▼ Use the UI

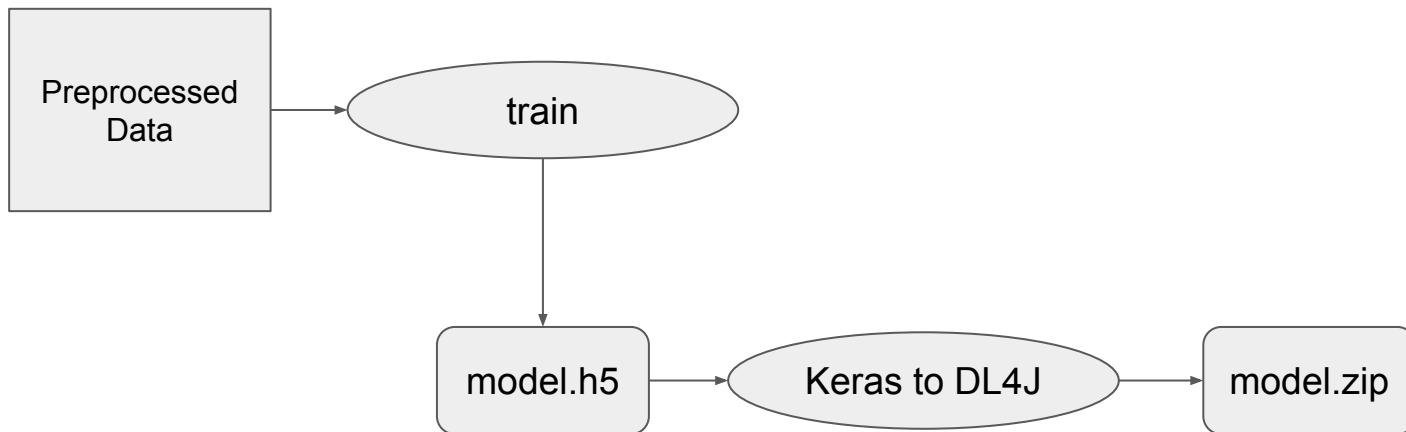


| Keras to DL4J

- ▼ Data Scientist loves Keras
- ▼ Keras is easier to import on notebook
- ▼ Training on Keras is faster
- ▼ Keras is compliant with cloud training (Sagemaker, CloudML)


```
def execute(config: Config): Unit = {  
  val kerasModel = KerasModelImport.importKerasModelAndWeights(  
    config.kerasModelPath, false)  
  ModelSaver.writeModel(kerasModel, config.outputModelPath)  
}
```


Workflow Train



Monitoring

Monitoring: mlflow


[Github](#) [Docs](#)

Listing Price Prediction

Experiment ID: 0 Artifact Location: /Users/matei/mlflow/demo/mlruns/0

Search Runs: [Search](#)

Filter Params: Filter Metrics: [Clear](#)

4 matching runs [Compare Selected](#) [Download CSV](#) 

	Time	User	Source	Version	Parameters		Metrics		
					alpha	l1_ratio	MAE	R2	RMSE
<input type="checkbox"/>	17:37	matei	linear.py	3a1995	0.5	0.2	84.27	0.277	158.1
<input type="checkbox"/>	17:37	matei	linear.py	3a1995	0.2	0.5	84.08	0.264	159.6
<input type="checkbox"/>	17:37	matei	linear.py	3a1995	0.5	0.5	84.12	0.272	158.6
<input type="checkbox"/>	17:37	matei	linear.py	3a1995	0	0	84.49	0.249	161.2

| Monitoring: mlflow

- ▼ Ensure your training machine can reach mlflow server
- ▼ Keep track of your experiment
 - ▽ Training parameter
 - ▽ Performance
- ▼ Compare results
- ▼ (model repository, standardize model packaging, easy deployment)



Monitoring: Zeppelin



Notebook ▾ Job


Search your Notes




anonymous ▾

ASTRO

MAE, RMSE



real vs prediction



| Monitoring: Zeppelin

- ▼ Already in HDP
- ▼ Authentication
- ▼ Scheduling
- ▼ Report View
- ▼ Auto shutdown
- ▼ Can mix sources (Scala, JDBC, C*, ...)
- ▼ API to automate deployment



Apache Zeppelin

Thank you for your attention

Any questions?