

Spark Deep Dive


Partition & Shuffle



<https://github.com/scauglog/prez>


| Vertical Approach

- ▼ Faster CPU
- ▼ More memory
- ▼ Simpler Programing
- ▼ Hardware Bounded
- ▼ Low volume




| Horizontal Approach





- ▼ Big Volume
- ▼ Many Server
- ▼ Complex programming
 - ▼ Failure Handling
 - ▼ Distributed Computing






Map Reduce






How many circle
and square?

















filter







map

















Group By













Group By















Group By








Group By →











Group By




Spark

Partitioning






PARTITION



PARTITION



PARTITION

Initial Partitioning

- ▼ depends on total size
- ▼ depends on parallelism
- ▼ small file -> 1 partition per 4 Mo
- ▼ intermediate -> 1 partition per task in parallel
- ▼ large file -> 1 partition per 128 Mo




```
Math.min(  
    maxPartitionBytes,  
    Math.max(openCostInBytes, bytesPerCore)  
)
```

```
spark.sql.files.maxPartitionBytes = 128 Mo
```




```
spark.sql.files.openCostInBytes = 4Mo
```

```
bytesPerCore = Total size / number of task in parallel
```




Parallelize
















filter








map

















TASK




TASK






map



filter



TASK






| Executor parallelism




- ▼ one executor process many partition in parallel
 - ▼ $\text{task in parallel} = \text{spark.executor.cores} / \text{spark.task.cpus}$
- ▼ limit the number of executor
 - ▼ less data send in case of broadcast
 - ▼ executor need some extra memory to work
- ▼ limit the number of cores by executor
 - ▼ executor with lots of core can't be allocated if nodes are too busy
 - ▼ performance drawback when reading file

Shuffle

everyday I'm shuffling



Group By



| Shuffle

- ▼ put all key on the same node
- ▼ create File in local FS (spark.shuffle.spill)
- ▼ different kind of shuffle
 - ▼ hash shuffle
 - ▼ consolidate hash shuffle
 - ▼ sort shuffle
 - ▼ tungsten sort shuffle

Partitioning RDD

- ▼ each aggregation operation has a signature with numPartitions
- ▼ `spark.default.parallelism`
 - ▽ ignored when working with dataframe
 - ▽ largest number of partition in a parent RDD
- ▼ `.coalesce()`
 - ▽ can only reduce number of partition
 - ▽ no shuffle involve
 - ▽ union of partition
- ▼ `.repartition()`
 - ▽ shuffle data
 - ▽ solve skewed data issue

```
myRdd.reduceByKey(_+_ , numPartitions = 14)
myRdd.sortByKey(true, numPartitions = 14)
myRdd.join(myRdd2, numPartitions = 14)
myRdd.coalesce(numPartitions = 5)
myRdd.repartition(numPartitions = 5)
```

| Partitioning DataFrame

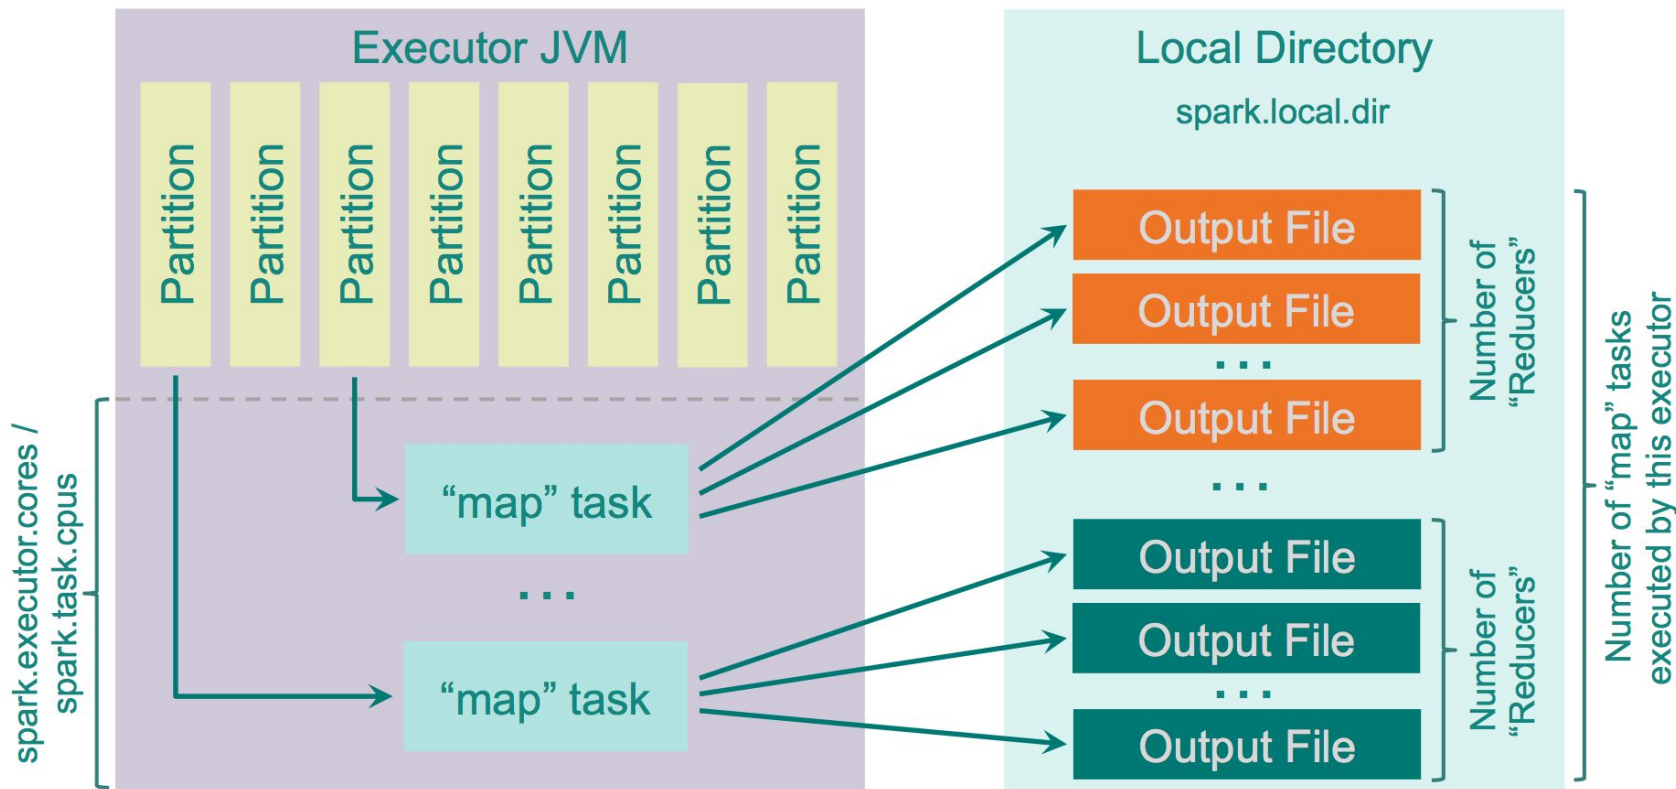
- ▼ we can't set a different numPartitions for shuffle operation
- ▼ `spark.sql.shuffle.partitions`
 - ▼ global for each shuffle operation
 - ▼ should be a multiple of the number of partition computed in parallel ($\text{num-executor} * \text{executor-cores} / \text{task.cores}$)
 - ▼ default = 200

| Hash Shuffle

- ▼ `spark.shuffle.manager=hash`
- ▼ 1 file for each reducer for each mapper
- ▼ number of file = $M * R$
- ▼ fast
- ▼ big amount of files written to FS
- ▼ lot of random IO

M: mapper task
R: reducer task

Hash Shuffle

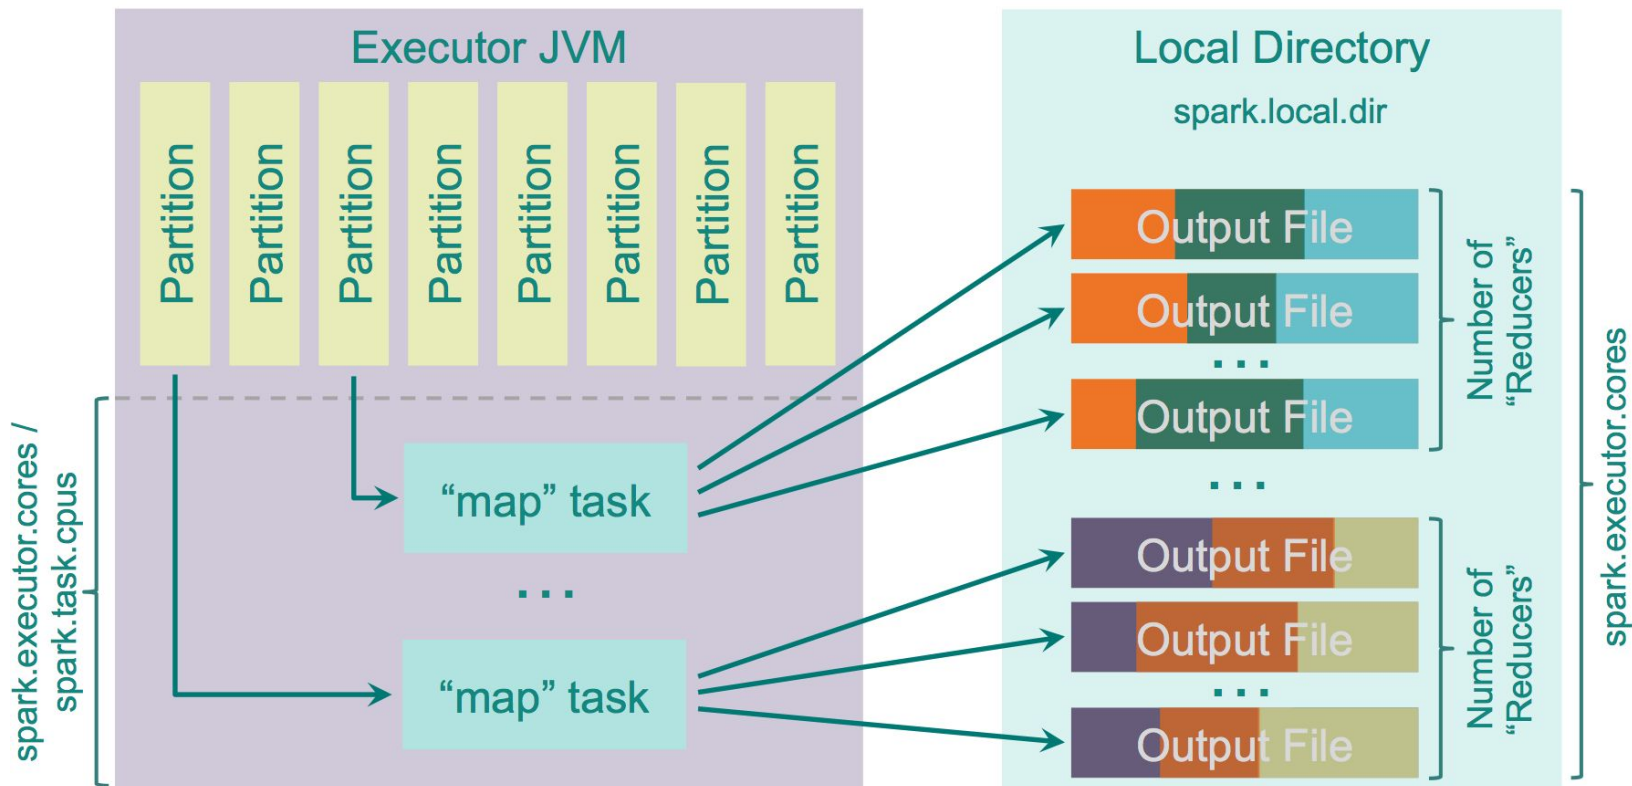


| Consolidate Hash Shuffle

- ▼ `spark.shuffle.manager=hash`
- ▼ `spark.shuffle.consolidateFiles=true`
- ▼ 1 file foreach reducer for each task in parallel by executor
- ▼ number of files = $E * C/T * R$
- ▼ less file written

M: mapper task
R: reducer task
E: num-executor
C: executor-cores
T: tasks.cpu

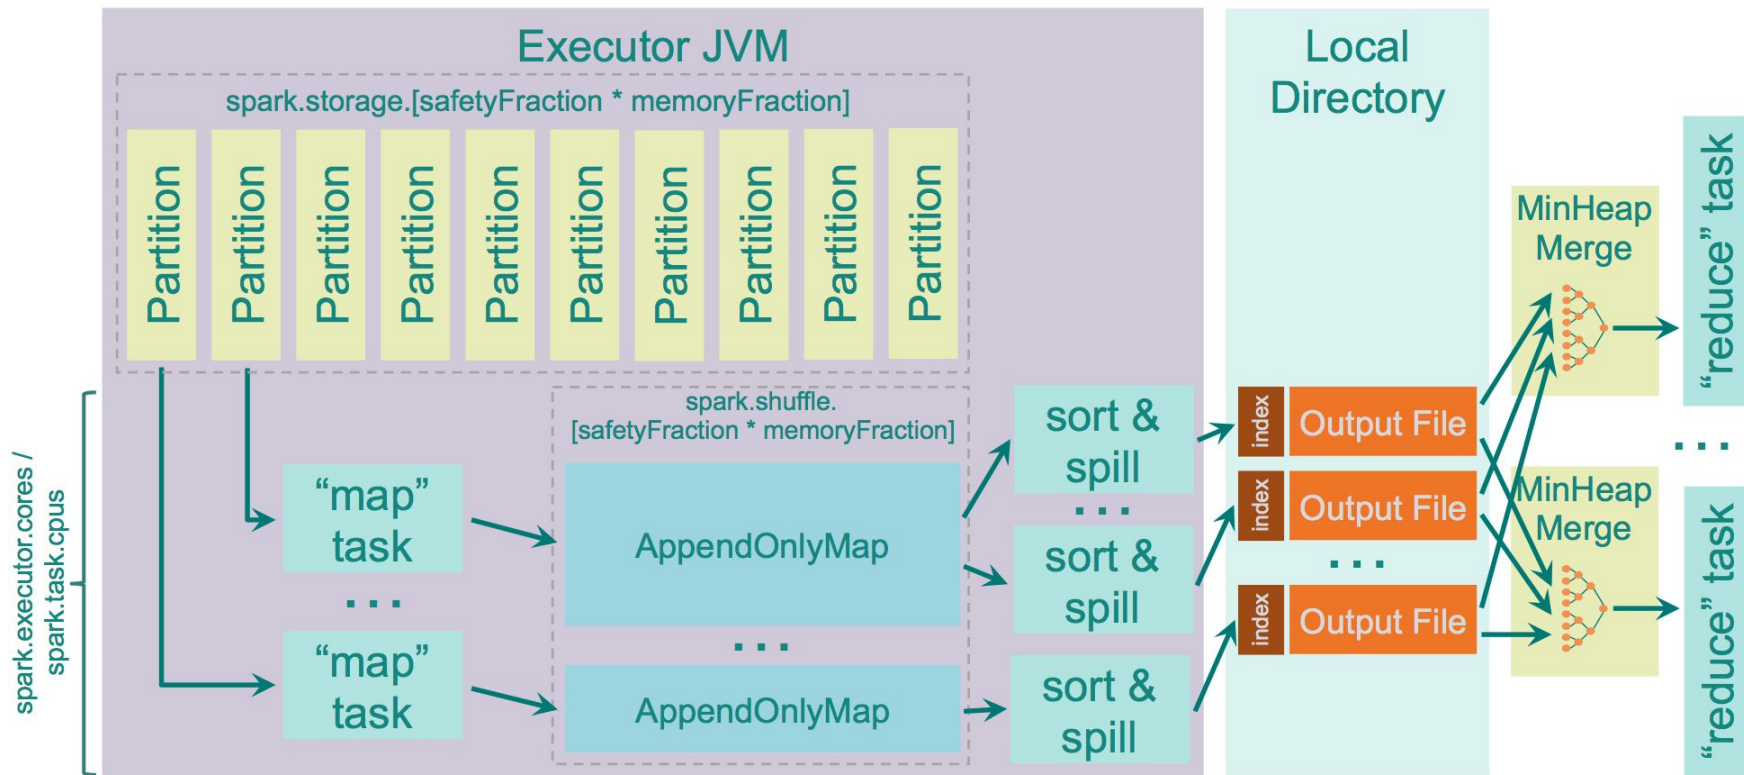
Consolidate Hash Shuffle



| Sort Shuffle

- ▼ `spark.shuffle.manager=sort`
- ▼ 1 file by mapper ordered by reducer and indexed
- ▼ if $R < 200$ then hash (`spark.shuffle.sort.bypassMergeThreshold`)
- ▼ sort data on map side using TimSort
- ▼ merge by reducer before sending to reducer
- ▼ sort after shuffle is faster

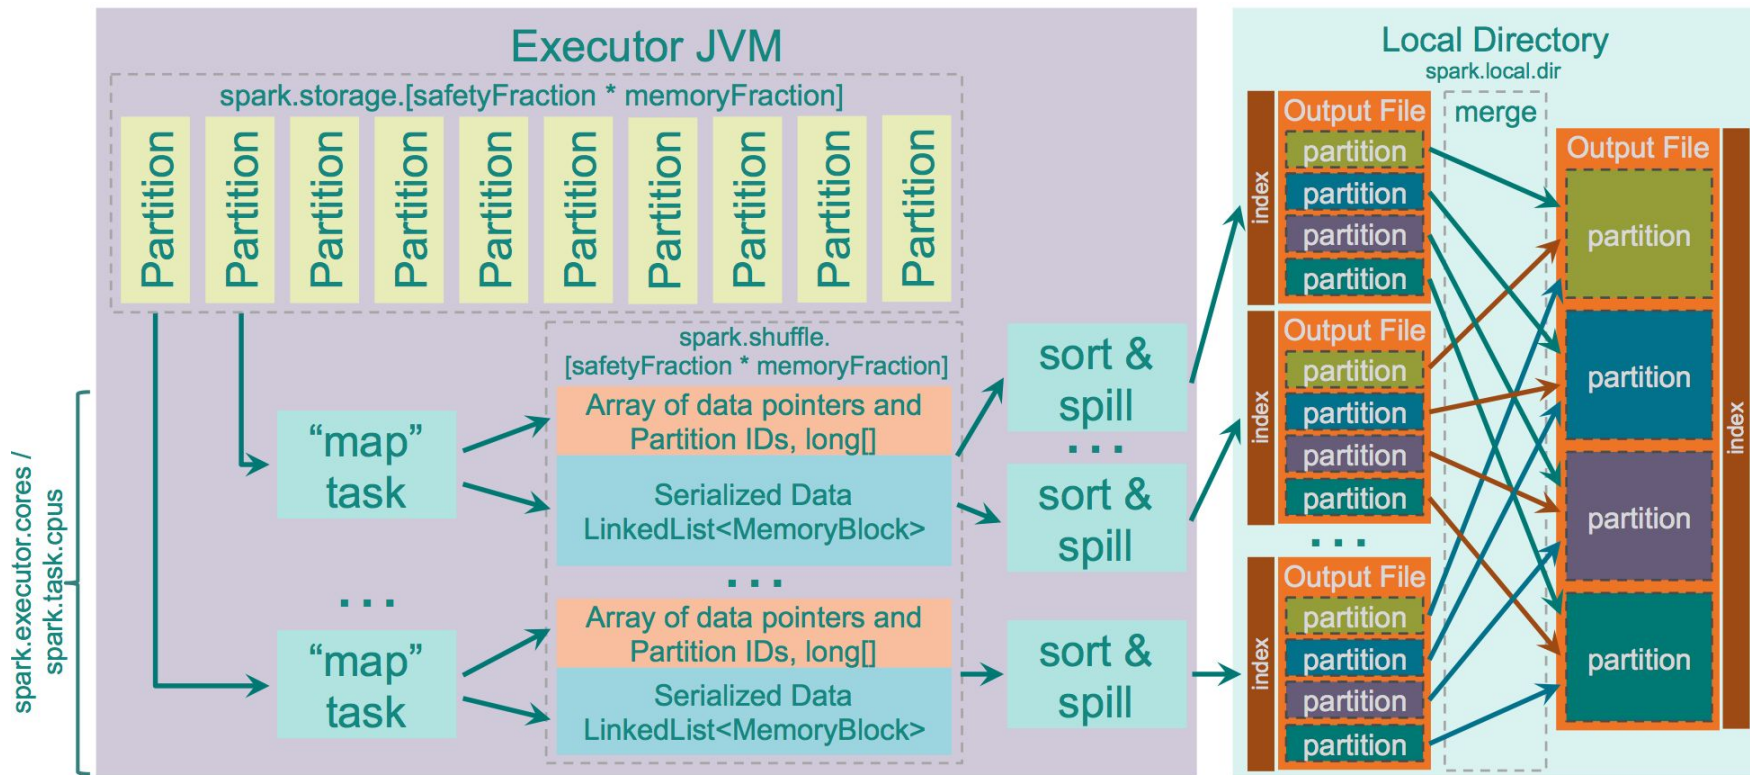
Sort Shuffle



| Tungsten Sort Shuffle

- ▼ `spark.shuffle.manager=tungsten-sort`
- ▼ operate on serialized data
- ▼ cache-efficient sorter
- ▼ work only if:
 - ▼ no aggregation (deserialisation)
 - ▼ less than 16 777 216 output partition (2^{24} , 3 octets)
 - ▼ row size < 128MB in serialized form
- ▼ no more fast sort after shuffle

Tungsten Sort Shuffle




| Shuffle take away



- ▼ hash shuffle
 - ▼ for small data
 - ▼ generate lots of file
- ▼ sort shuffle
 - ▼ for large data
 - ▼ slower due to sort
 - ▼ less file generated
- ▼ tune with `spark.shuffle.sort.bypassMergeThreshold`

JOIN

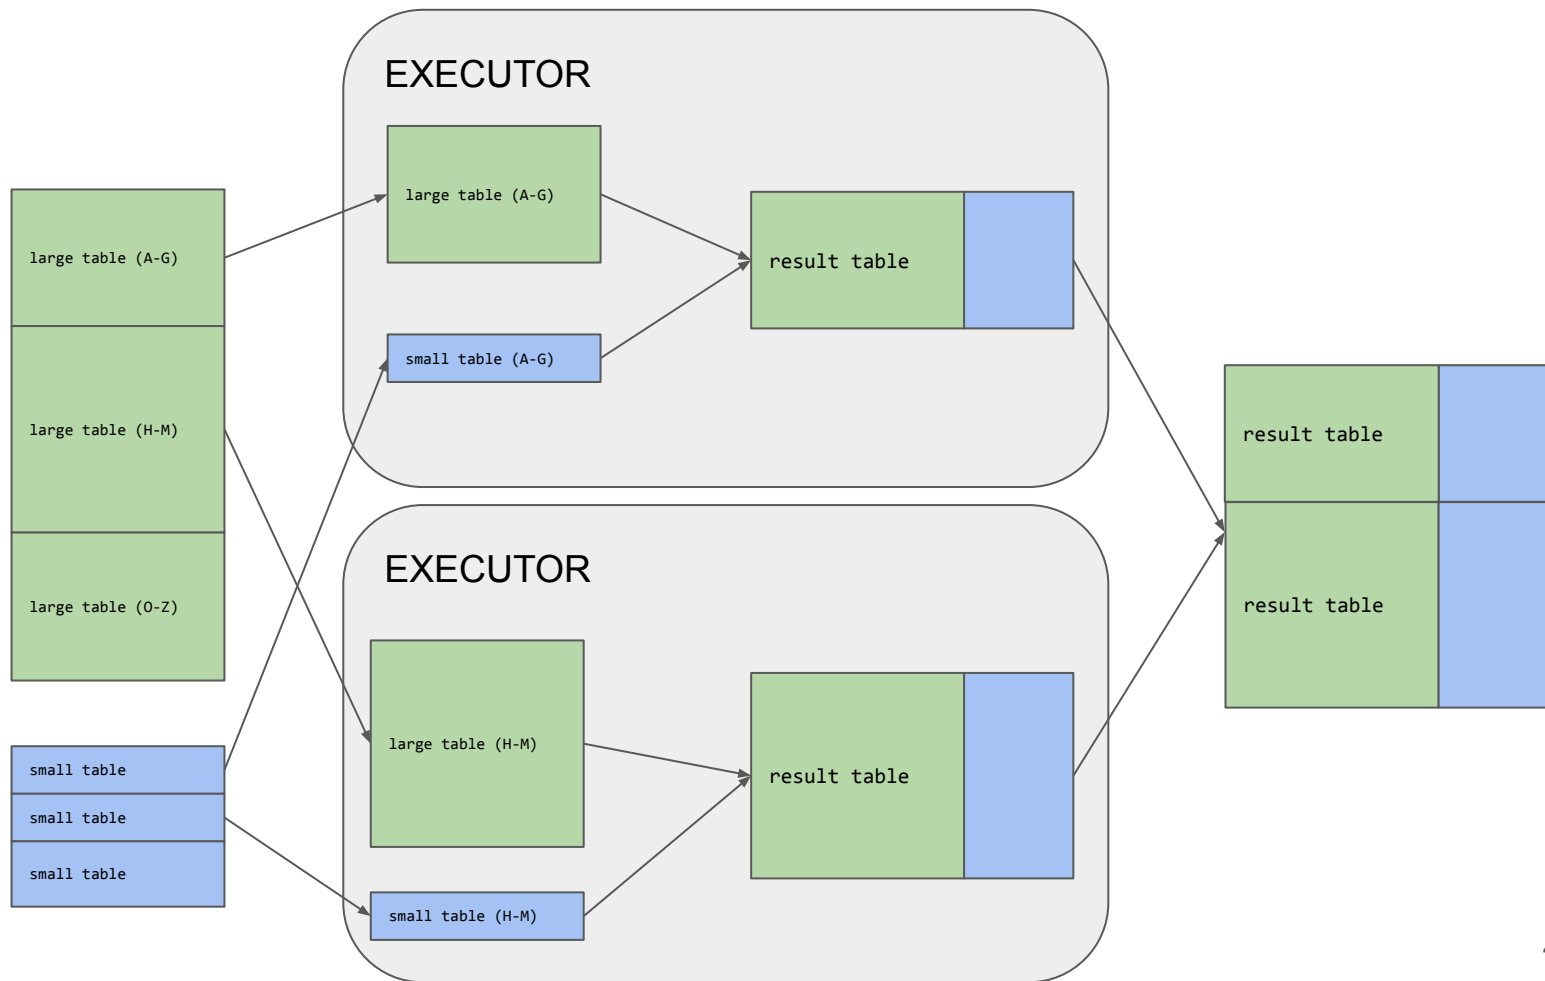
Join



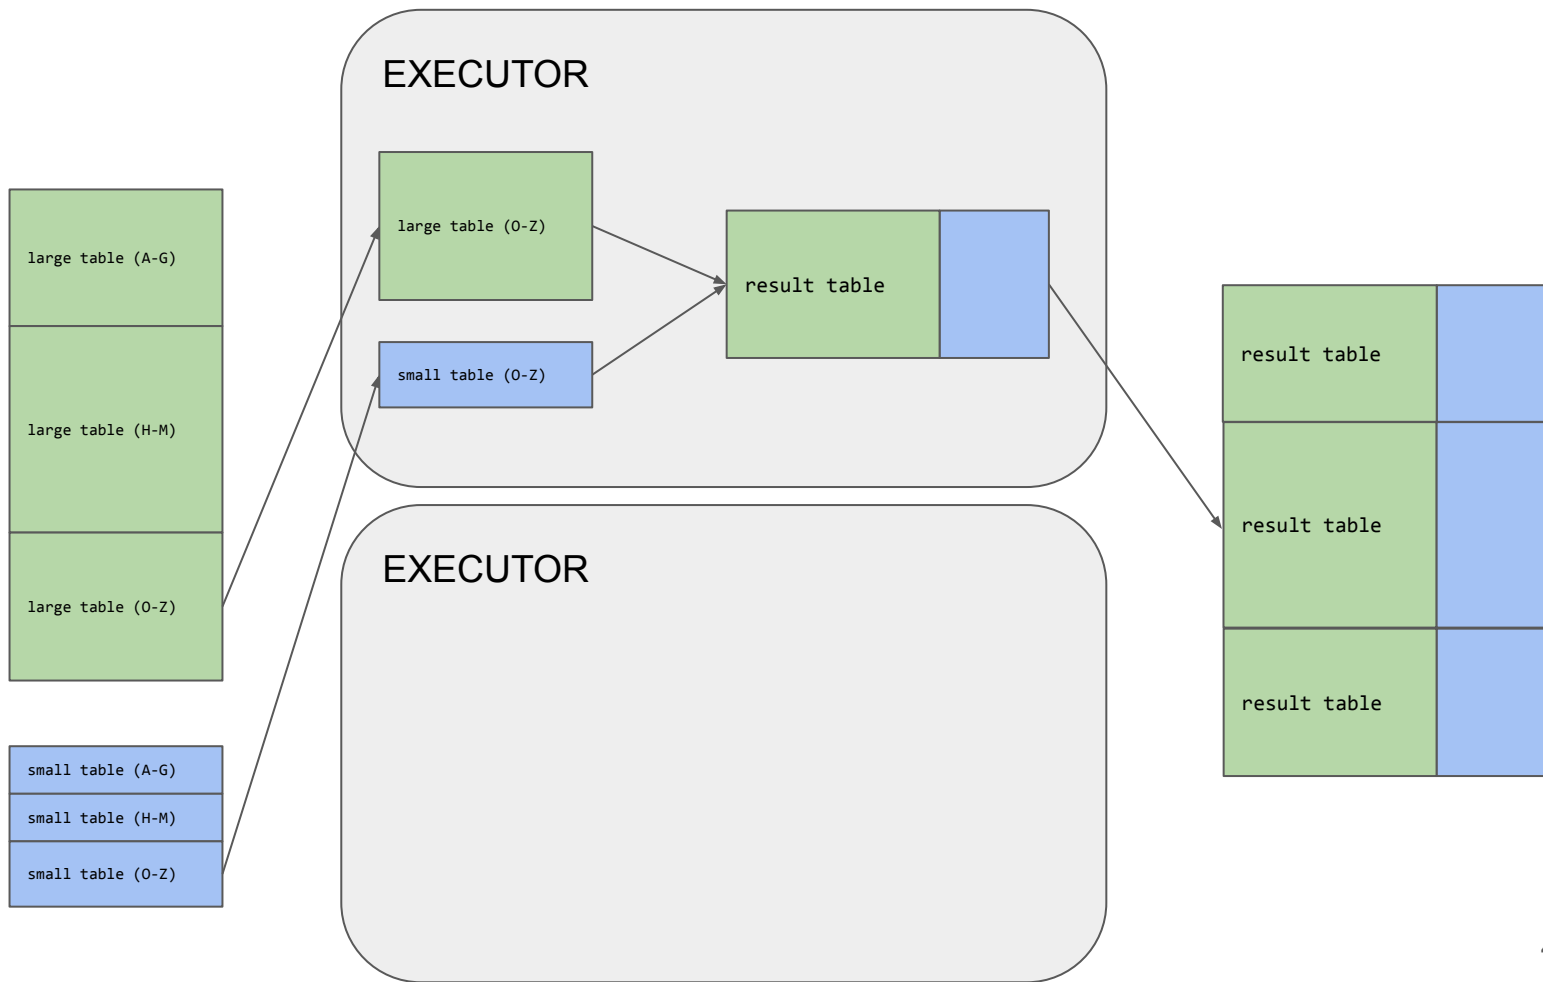
Join





Join



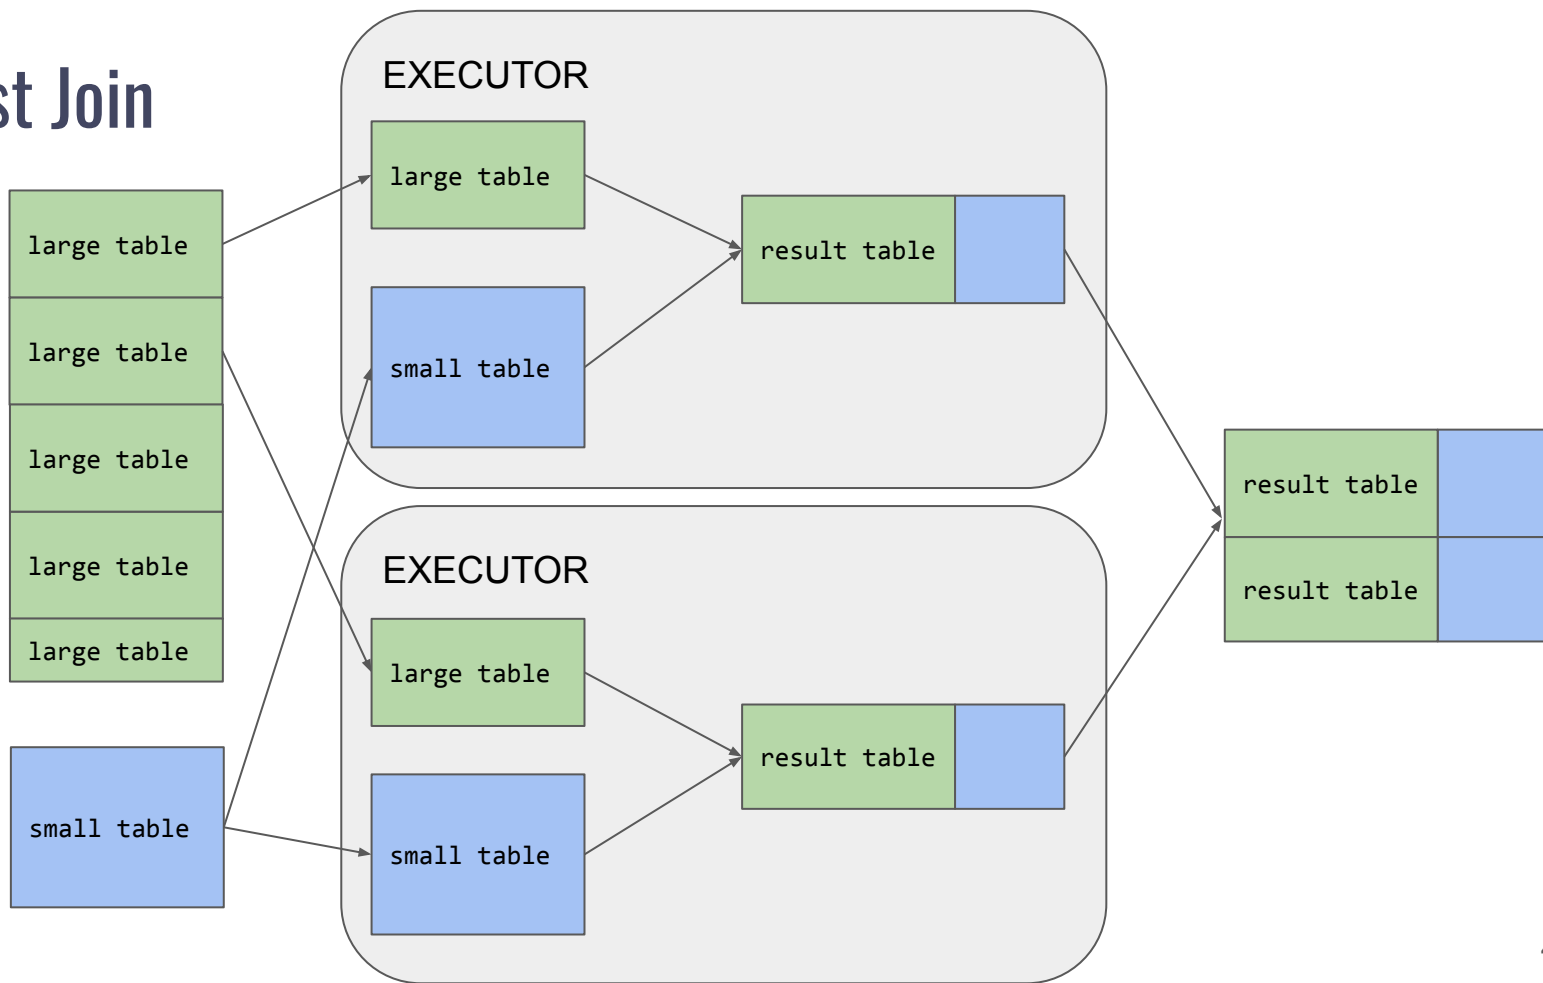
Join



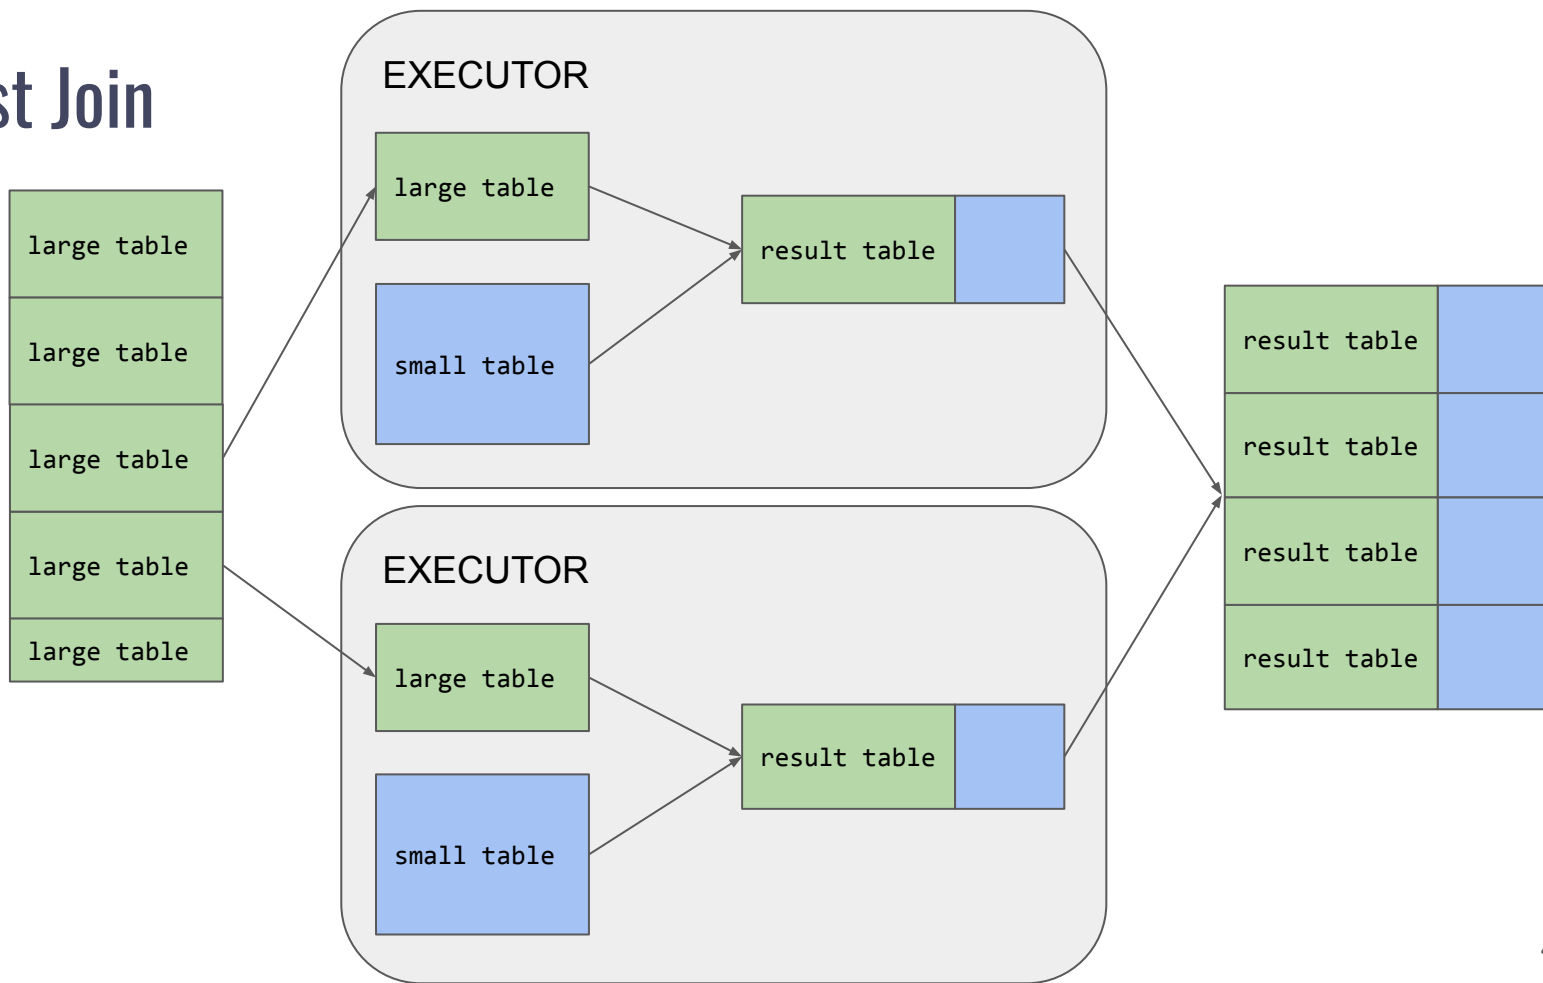
Broadcast Join



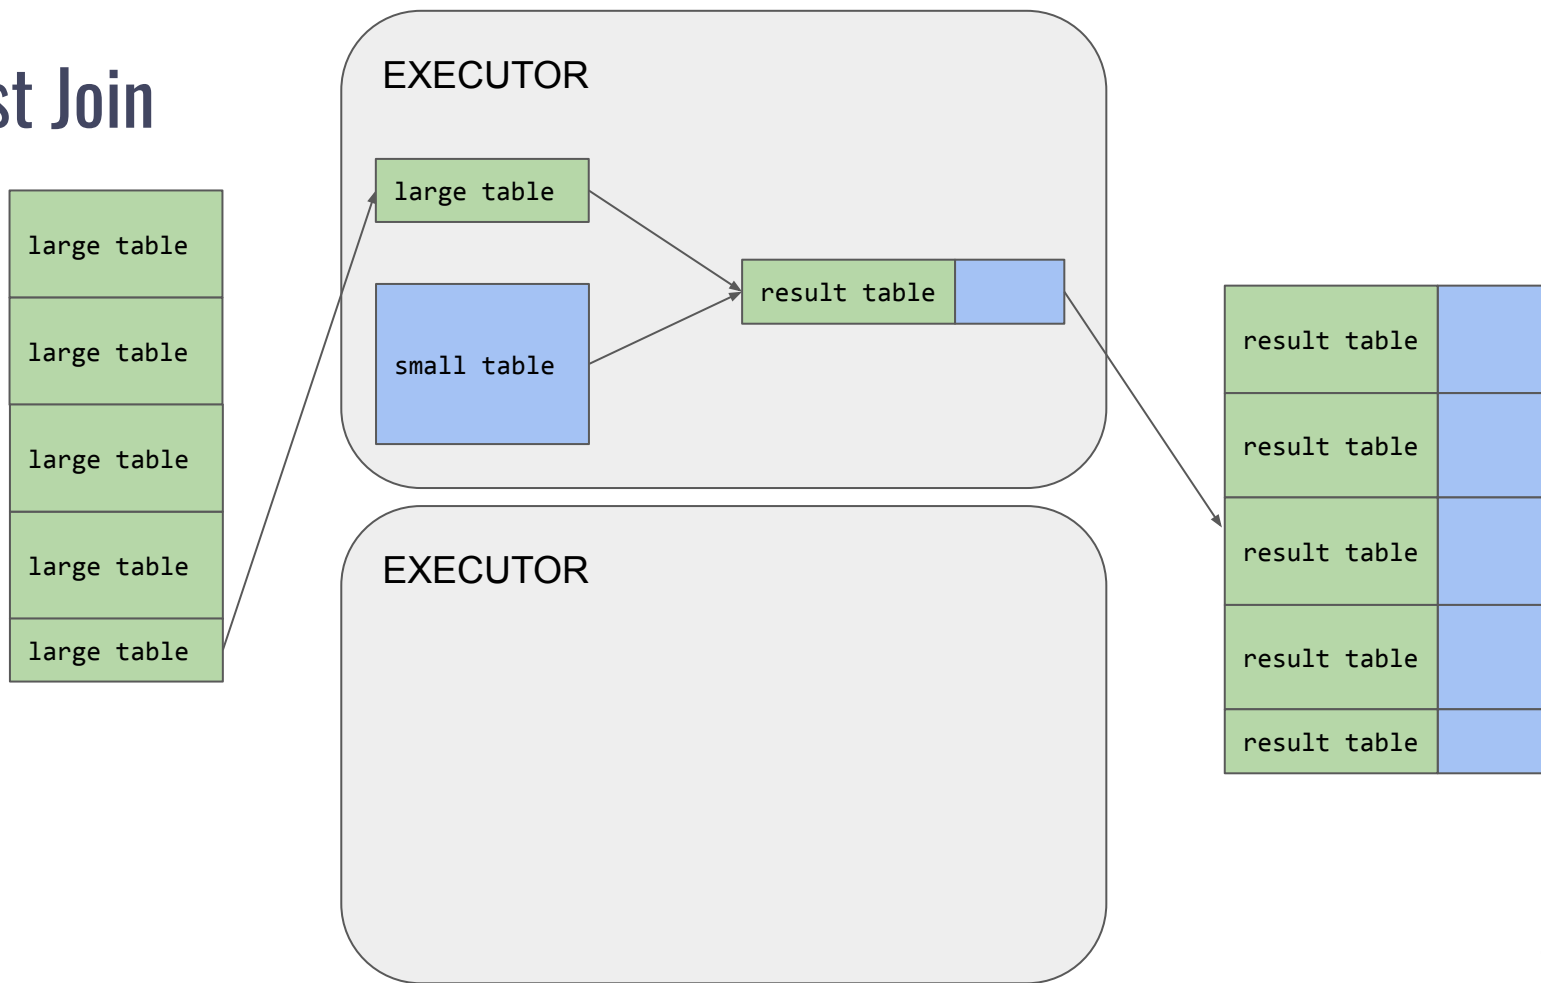
Broadcast Join



Broadcast Join



Broadcast Join



| Broadcast join


- ▼ no shuffle
- ▼ the small table is sent (broadcast) to all executor
- ▼ the small table must fit in executor memory
- ▼ broadcast join is automatic but it's better to specify explicitly
- ▼ `spark.sql.autoBroadcastJoinThreshold = 10Mb`

```
df1.join(broadcast(df2), Seq("column1"))
```

Tasks, Stages & Jobs


| Task and Stage

```
val wc = sc.textFile(myData)
```



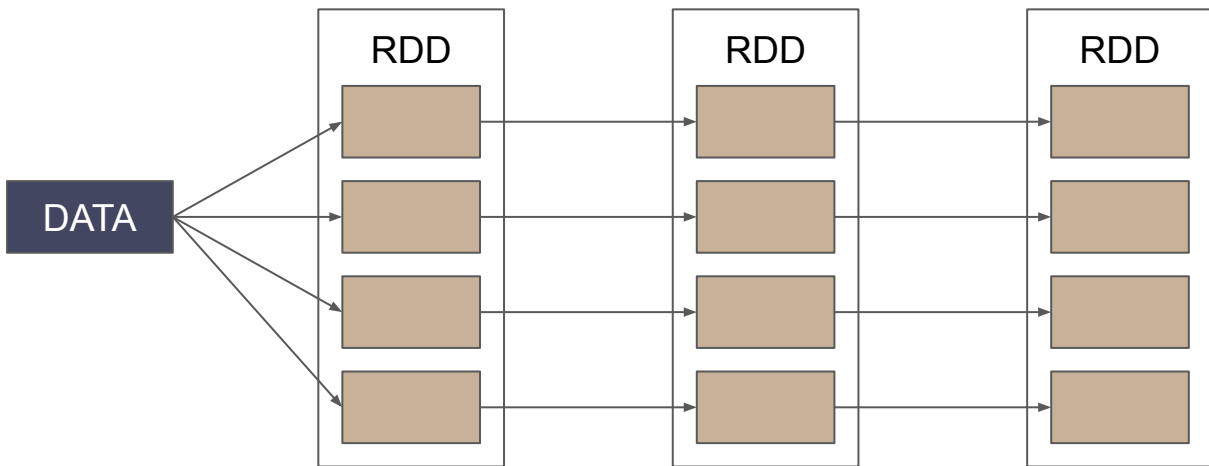
| Task and Stage

```
val wc = sc.textFile(myData)  
    .map(_._split(","))
```




| Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
```



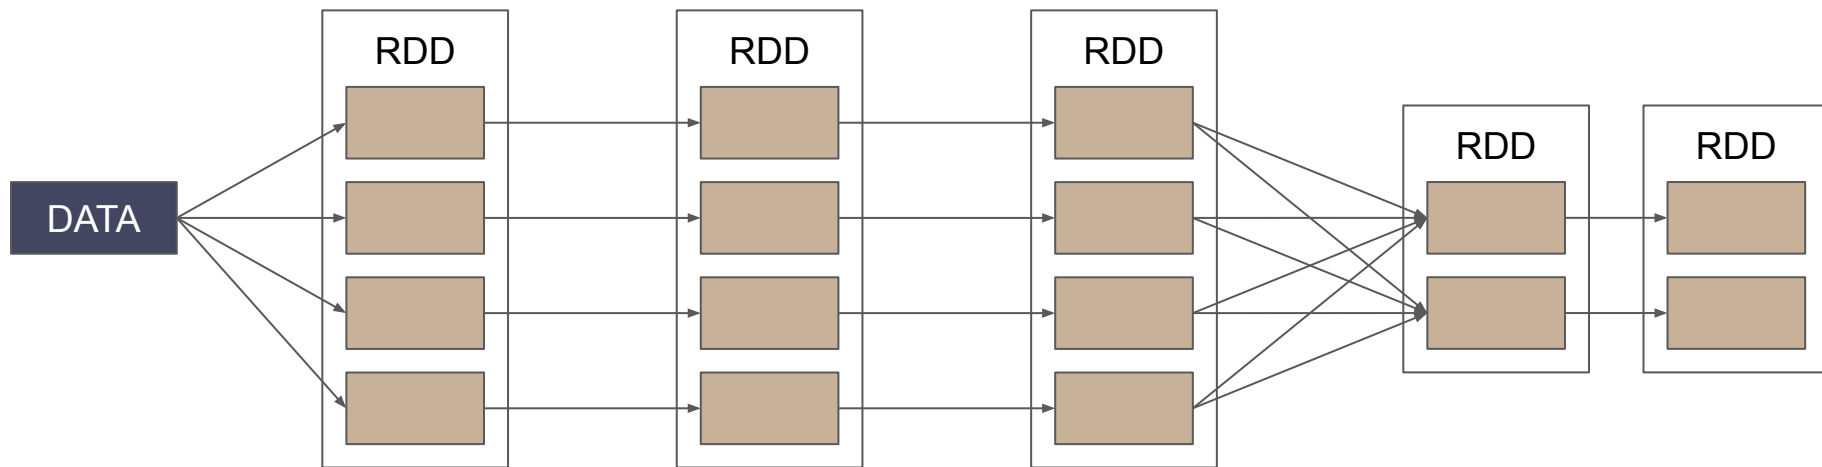
Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
```




Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
  .mapValues { case (dataSum, cpt) => dataSum/cpt }
```




Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
  .mapValues { case (dataSum, cpt) => dataSum/cpt }
```




Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
  .mapValues { case (dataSum, cpt) => dataSum/cpt }
```



Task and Stage


```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
  .mapValues { case (dataSum, cpt) => dataSum/cpt }
```




Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
  .mapValues { case (dataSum, cpt) => dataSum/cpt }
```

STAGE 0




STAGE 1



Task and Stage

```
val wc = sc.textFile(myData)
  .map(_.split(","))
  .map(row => (row(0), (row(14), 1)))
  .reduceByKey( (x,y) => (x._1 + y._1, x._2 + y._2))
  .mapValues { case (dataSum, cpt) => dataSum/cpt }
```



Spark UI: Jobs

[Jobs](#)[Stages](#)[Storage](#)[Environment](#)[Executors](#)[SQL](#)[Zeppelin application UI](#)

Spark Jobs (?)

User: sagean

Total Uptime: 5,2 min

Scheduling Mode: FIFO

Completed Jobs: 1

[▶ Event Timeline](#)

Completed Jobs (1)

Job Id (Job Group) ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0 (zeppelin-2E18ZAS6F-20181228-173254_1752340159)	Started by: admin saveAsTextFile at <console>:28	2019/01/21 16:09:03	1 s	2/2	4/4

Spark UI: Stage

[Jobs](#)[Stages](#)[Storage](#)[Environment](#)[Executors](#)[SQL](#)

Zeppelin application UI

Details for Job 0

Status: SUCCEEDED**Job Group:** zeppelin-2E18ZAS6F-20181228-173254_1752340159**Completed Stages:** 2[▶ Event Timeline](#)[▶ DAG Visualization](#)

Completed Stages (2)

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	Started by: admin saveAsTextFile at <console>:28 +details	2019/01/21 16:09:05	0,2 s	<div>2/2</div>		11.3 KB	11.3 KB	
0	Started by: admin map at <console>:25 +details	2019/01/21 16:09:04	1,0 s	<div>2/2</div>	14.7 KB			11.3 KB

Spark UI: Stage DAG


Details for Job 0

Status: SUCCEEDED

Job Group: zeppelin-2E18ZAS6F-20181228-173254_1752340159

Completed Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization



Spark UI: Tasks

JobsStagesStorageEnvironmentExecutorsSQLZeppelin application UI

Details for Stage 1 (Attempt 0)

Total Time Across All Tasks: 0,3 s

Locality Level Summary: Any: 2

Output: 11.3 KB / 957

Shuffle Read: 11.3 KB / 977

- ▶ DAG Visualization
- ▶ Show Additional Metrics
- ▶ Event Timeline

Summary Metrics for 2 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0,1 s	0,1 s	0,2 s	0,2 s	0,2 s
GC Time	0 ms	0 ms	0 ms	0 ms	0 ms
Output Size / Records	5.5 KB / 462	5.5 KB / 462	5.8 KB / 495	5.8 KB / 495	5.8 KB / 495
Shuffle Read Size / Records	5.5 KB / 470	5.5 KB / 470	5.8 KB / 507	5.8 KB / 507	5.8 KB / 507

▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks	Output Size / Records	Shuffle Read Size / Records
driver	10.7.14.164:36674	0,4 s	2	0	0	2	11.3 KB / 957	11.3 KB / 977

Tasks (2)

Index ▲	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Output Size / Records	Shuffle Read Size / Records	Errors
0	2	0	SUCCESS	ANY	driver / localhost	2019/01/21 16:09:05	0,1 s		5.8 KB / 495	5.8 KB / 507	
1	3	0	SUCCESS	ANY	driver / localhost	2019/01/21 16:09:05	0,2 s		5.5 KB / 462	5.5 KB / 470	

The End